

# Activity\_Course 2 TikTok project lab

August 26, 2023

## 1 TikTok Project

**Goals:** 1. Acquaint you with the data

2. Compile summary information about the data
3. Begin the process of EDA and reveal insights contained in the data
4. Prepare you for more in-depth EDA, hypothesis testing, and statistical analysis

**Part 1:** Understand the situation \* How can you best prepare to understand and organize the provided TikTok information?

**Part 2:** Understand the data

- Create a pandas dataframe for data learning and future exploratory data analysis (EDA) and statistical activities
- Compile summary information about the data to inform next steps

**Part 3:** Understand the variables

- Use insights from your examination of the summary data to guide deeper investigation into variables

### 1.0.1 Imports and data loading

```
[1]: import pandas as pd
```

```
import numpy as np
```

```
[2]: # Load dataset into dataframe  
data = pd.read_csv("tiktok_dataset.csv")
```

### 1.0.2 Understand the data - Inspect the data

**Question 1:** When reviewing the first few rows of the dataframe, what do you observe about the data? What does each row represent?

**Question 2:** When reviewing the `data.info()` output, what do you notice about the different variables? Are there any null values? Are all of the variables numeric? Does anything else stand out?

**Question 3:** When reviewing the `data.describe()` output, what do you notice about the distributions of each variable? Are there any questionable values? Does it seem that there are outlier values?

```
[3]: data.head(10)
#Each row represents the claim number, its status, an associated video ID and
→its length, the transcript text of the video, its verified and banned
→status, and its view, like, share, download, and comment counts.
```

```
[3]:
```

	#	claim_status	video_id	video_duration_sec	\
0	1	claim	7017666017	59	
1	2	claim	4014381136	32	
2	3	claim	9859838091	31	
3	4	claim	1866847991	25	
4	5	claim	7105231098	19	
5	6	claim	8972200955	35	
6	7	claim	4958886992	16	
7	8	claim	2270982263	41	
8	9	claim	5235769692	50	
9	10	claim	4660861094	45	

  

	video_transcription_text	verified_status	\
0	someone shared with me that drone deliveries a...	not verified	
1	someone shared with me that there are more mic...	not verified	
2	someone shared with me that american industria...	not verified	
3	someone shared with me that the metro of st. p...	not verified	
4	someone shared with me that the number of busi...	not verified	
5	someone shared with me that gross domestic pro...	not verified	
6	someone shared with me that elvis presley has ...	not verified	
7	someone shared with me that the best selling s...	not verified	
8	someone shared with me that about half of the ...	not verified	
9	someone shared with me that it would take a 50...	verified	

  

	author_ban_status	video_view_count	video_like_count	video_share_count	\
0	under review	343296.0	19425.0	241.0	
1	active	140877.0	77355.0	19034.0	
2	active	902185.0	97690.0	2858.0	
3	active	437506.0	239954.0	34812.0	
4	active	56167.0	34987.0	4110.0	
5	under review	336647.0	175546.0	62303.0	
6	active	750345.0	486192.0	193911.0	
7	active	547532.0	1072.0	50.0	
8	active	24819.0	10160.0	1050.0	
9	active	931587.0	171051.0	67739.0	

	video_download_count	video_comment_count
0	1.0	0.0
1	1161.0	684.0
2	833.0	329.0
3	1234.0	584.0
4	547.0	152.0
5	4293.0	1857.0
6	8616.0	5446.0
7	22.0	11.0
8	53.0	27.0
9	4104.0	2540.0

```
[4]: data.info()
#There are 298 rows with missing claim status, and video view, like, share,
↳download, and comment data. This represents 1.5% of the data so it might be
↳best during the cleaning stage to remove those rows.
#View, like, share, download, and comment are Float values rather than Int
↳values even though those values are clearly integers.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19382 entries, 0 to 19381
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   #                                       19382 non-null  int64
1   claim_status                          19084 non-null  object
2   video_id                              19382 non-null  int64
3   video_duration_sec                    19382 non-null  int64
4   video_transcription_text              19084 non-null  object
5   verified_status                       19382 non-null  object
6   author_ban_status                     19382 non-null  object
7   video_view_count                      19084 non-null  float64
8   video_like_count                      19084 non-null  float64
9   video_share_count                     19084 non-null  float64
10  video_download_count                  19084 non-null  float64
11  video_comment_count                   19084 non-null  float64
dtypes: float64(5), int64(3), object(4)
memory usage: 1.8+ MB
```

```
[5]: data.describe()
#Mean video duration is 32.4 seconds, the STDV is high at 16.2
#The shortest video is 5 seconds.
#The longest is 60 (this makes sense because TikTok limited videos to 60
↳seconds at this time, so it might be interesting to compare the distribution
↳of video lengths).
```

```
#Mean views, likes, shares, downloads, and comments are all lower than their
↳STDVs, indicating that there are some extreme upper end outliers to be
↳investigated. Perhaps there is a way to divide the data into subsets based
↳on their view counts.
#This is reinforced by the fact that many of the max values are 100x larger
↳than their medians
```

```
[5]:
```

	#	video_id	video_duration_sec	video_view_count	\
count	19382.000000	1.938200e+04	19382.000000	19084.000000	
mean	9691.500000	5.627454e+09	32.421732	254708.558688	
std	5595.245794	2.536440e+09	16.229967	322893.280814	
min	1.000000	1.234959e+09	5.000000	20.000000	
25%	4846.250000	3.430417e+09	18.000000	4942.500000	
50%	9691.500000	5.618664e+09	32.000000	9954.500000	
75%	14536.750000	7.843960e+09	47.000000	504327.000000	
max	19382.000000	9.999873e+09	60.000000	999817.000000	

  

	video_like_count	video_share_count	video_download_count	\
count	19084.000000	19084.000000	19084.000000	
mean	84304.636030	16735.248323	1049.429627	
std	133420.546814	32036.174350	2004.299894	
min	0.000000	0.000000	0.000000	
25%	810.750000	115.000000	7.000000	
50%	3403.500000	717.000000	46.000000	
75%	125020.000000	18222.000000	1156.250000	
max	657830.000000	256130.000000	14994.000000	

  

	video_comment_count
count	19084.000000
mean	349.312146
std	799.638865
min	0.000000
25%	1.000000
50%	9.000000
75%	292.000000
max	9599.000000

### 1.0.3 Understand the data - Investigate the variables

The ultimate objective is to use machine learning to classify videos as either claims or opinions. Examine the `claim_status` variable by determining how many videos there are for each different claim status:

```
[6]: pd.Series(data['claim_status']).value_counts()
#The number of claims and opinions are similar and there are 298 null values as
↳previously noted.
```

```
[6]: claim      9608
      opinion     9476
      Name: claim_status, dtype: int64
```

Using a Boolean masking to filter the data according to claim status, then calculate the mean and median view counts for each claim status:

```
[7]: # What is the average view count of videos with "claim" status?
      claim_mask = data[data.claim_status == 'claim']
      claim_av = np.average(claim_mask['video_view_count'])
      claim_med = np.median(claim_mask['video_view_count'])
      print('Claim views mean:', claim_av, 'Claim views median:', claim_med, sep="\n")
      #The average and median are very similar. This indicates that the views counts
      ↳ of videos with claims are likely to be evenly distributed.
```

```
Claim views mean:
501029.4527477102
Claim views median:
501555.0
```

```
[8]: # What is the average view count of videos with "opinion" status?
      op_mask = data[data.claim_status == 'opinion']
      op_av = np.average(op_mask['video_view_count'])
      op_med = np.median(op_mask['video_view_count'])
      print('Opinion views mean:', op_av, 'Opinion views median:', op_med, sep="\n")
      #The average and median are very similar. This indicates that the views counts
      ↳ of videos with claims are likely to be evenly distributed.
      #Claim videos receive an average of 100x as many views as opinion videos.
```

```
Opinion views mean:
4956.43224989447
Opinion views median:
4953.0
```

Use `groupby()` to calculate how many videos there are for each combination of categories of claim status and author ban status.

```
[9]: # Get counts for each group combination of claim status and author ban status
      data.groupby(['claim_status', 'author_ban_status']).size()
      #Banned authors who make claims have a much higher percentage of their category
      ↳ than banned others who present opinions.
      #Authors who make claims are more likely to make claims that are innacurate and
      ↳ claims that are innacurate are more likely to result in a ban.
```

```
[9]: claim_status  author_ban_status
      claim        active          6566
           banned          1439
           under review    1603
      opinion        active          8817
```

```

        banned          196
        under review    463
dtype: int64

```

Calculate the median video share count of each author ban status.

```

[10]: ban_mask = data[data.author_ban_status == 'banned']
      active_mask = data[data.author_ban_status == 'active']
      ur_mask = data[data.author_ban_status == 'under review']

      ban_med = np.nanmedian(ban_mask['video_share_count'])
      active_med = np.nanmedian(active_mask['video_share_count'])
      ur_med = np.nanmedian(ur_mask['video_share_count'])

      print('Banned median:',ban_med,'Active median:',active_med,'Under review median:
      ↪',ur_med, sep="\n")

```

```

Banned median:
14468.0
Active median:
437.0
Under review median:
9444.0

```

```

[11]: # What's the median video share count of each author ban status?
      #Banned: 14468
      #Active: 437
      #Under review: 9444
      #The median share count of banned authors is massive in comparison to active_
      ↪authors.

```

Use `groupby()` to group the data by `author_ban_status`, then use `agg()` to get the count, mean, and median of each of the following columns: \* `video_view_count` \* `video_like_count` \* `video_share_count`

```

[12]: ban_views = data.groupby(['author_ban_status']).
      ↪agg(TotalViews=('video_view_count', 'sum'), MeanViews=('video_view_count', 'mean'), MedViews=('v

      ban_likes = data.groupby(['author_ban_status']).
      ↪agg(TotalLikes=('video_like_count', 'sum'), MeanLikes=('video_like_count', 'mean'), MedLikes=('v

      ban_shares = data.groupby(['author_ban_status']).
      ↪agg(TotalShares=('video_share_count', 'sum'), MeanShares=('video_share_count', 'mean'), MedShar

      print(ban_views, ban_likes, ban_shares, sep='\n\n')
      #Active authors have about 50% more views and likes than banned authors and 5x_
      ↪more shares than banned authors.

```

*#Banned authors have significantly higher mean and median views, likes, and shares than active authors though.*  
*#Specifically, banned authors have massive mean and median shares compared to active authors.*  
*#Perhaps those who watch videos from banned authors are more likely to share videos.*

	TotalViews	MeanViews	MedViews
author_ban_status			
active	3.321606e+09	215927.039524	8616.0
banned	7.289573e+08	445845.439144	448201.0
under review	8.102952e+08	392204.836399	365245.5

	TotalLikes	MeanLikes	MedLikes
author_ban_status			
active	1.092755e+09	71036.533836	2222.0
banned	2.501832e+08	153017.236697	105573.0
under review	2.659315e+08	128718.050339	71204.5

	TotalShares	MeanShares	MedShares
author_ban_status			
active	217076684.0	14111.466164	437.0
banned	49048271.0	29998.942508	14468.0
under review	53250524.0	25774.696999	9444.0

Create three new columns to help better understand engagement rates: \* `likes_per_view`: represents the number of likes divided by the number of views for each video \* `comments_per_view`: represents the number of comments divided by the number of views for each video \* `shares_per_view`: represents the number of shares divided by the number of views for each video

```
[13]: # Create a likes_per_view column
data['LPV'] = data['video_like_count']/data['video_view_count']

# Create a comments_per_view column
data['CPV'] = data['video_comment_count']/data['video_view_count']

# Create a shares_per_view column
data['SPV'] = data['video_share_count']/data['video_view_count']
```

Use `groupby()` to compile the information in each of the three newly created columns for each combination of categories of claim status and author ban status, then use `agg()` to calculate the count, the mean, and the median of each group.

```
[14]: ban_LPV = data.groupby(['author_ban_status', 'claim_status']).
    →agg(TotalLPV=('LPV', 'sum'), MeanLPV=('LPV', 'mean'), MedLPV=('LPV', 'median'))
#LikesPerView
ban_CPV = data.groupby(['author_ban_status', 'claim_status']).
    →agg(TotalCPV=('CPV', 'sum'), MeanCPV=('CPV', 'mean'), MedCPV=('CPV', 'median'))
```

```

#CommentsPerView
ban_SPV = data.groupby(['author_ban_status', 'claim_status']).
    ↳agg(TotalSPV=('SPV', 'sum'), MeanSPV=('SPV', 'mean'), MedSPV=('SPV', 'median'))
#SharesPerView
print(ban_LPV, ban_CPV, ban_SPV, sep='\n\n')
#Likes, commends, and share are more common for claims than opinions, but over
    ↳10x as likely for banned users.
#It appears that authors who are banned got a much more significant portion of
    ↳their engagement from claims than from opinions.

```

		TotalLPV	MeanLPV	MedLPV
author_ban_status	claim_status			
active	claim	2163.772970	0.329542	0.326538
	opinion	1937.478628	0.219744	0.218330
banned	claim	496.556528	0.345071	0.358909
	opinion	40.546207	0.206868	0.198483
under review	claim	525.778643	0.327997	0.320867
	opinion	104.820594	0.226394	0.228051

		TotalCPV	MeanCPV	MedCPV
author_ban_status	claim_status			
active	claim	9.144001	0.001393	0.000776
	opinion	4.559119	0.000517	0.000252
banned	claim	1.981775	0.001377	0.000746
	opinion	0.085135	0.000434	0.000193
under review	claim	2.191442	0.001367	0.000789
	opinion	0.247962	0.000536	0.000293

		TotalSPV	MeanSPV	MedSPV
author_ban_status	claim_status			
active	claim	429.782704	0.065456	0.049279
	opinion	385.554840	0.043729	0.032405
banned	claim	97.698624	0.067893	0.051606
	opinion	7.944022	0.040531	0.030728
under review	claim	105.370751	0.065733	0.049967
	opinion	20.590731	0.044472	0.035027

Executive summary for this TikTok project

Here is an initial summary of the findings in the TikTok data set which could impact the ability to create a model for predicting whether a video contains a claim or an opinion.

There are 298 rows with missing claim status, and video view, like, share, download, and comment data. This represents 1.5% of the data so it might be best during the cleaning stage to remove those rows or impute 0s. Number, video ID, and video duration are integers even though video duration could be a float value, but it seems that video duration was only measured to the nearest second. It is not indicated what type of rounding was used. Claim, transcript, verified, and banned are all object values. Claim has 'claim' and 'opinion', verified has 'verified' and 'not verified', and banned has 'banned', 'active', and 'under review.' View, like, share, download, and comment are



Float values rather than Int values even though those values are clearly integers.

The mean views, likes, and shares of active authors are all many times larger than their medians, indicating a large amount of skew in the data for active authors. The mean likes and shares of banned authors are also 2-3x as large as their medians, indicating some skew in the data for banned authors. When broken down into a per view basis, likes are more likely to be normally distributed, comments are much more likely to be skewed, and shares are somewhat likely to be skewed. Skew measurements for each of these might provide more insight about the types of tests or analysis tools could be used to make accurate predictions given their varying levels of skew.

Video durations have a minimum of 5 seconds and a maximum of 60 as per TikToks previous video length restrictions. This might impact reliability of models built off this data since videos are now allowed to be longer. Video view counts range the most from 20 to almost 1 million with a heavy skew to the right Like, comment, download, and share values all range from 0 to 650k, 250k, 15k, and 10k respectively and all of these are heavily skewed to the right.

Specifically with respect to the variables in question: claims and opinions: The average and median views of videos that are claims are approximately 501k and 502k respectively. The average and median views of videos that are opinions are approximately 4956 and 4953 respectively.

I would recommend further exploration before beginning EDA. A lot of insight could be gained about the way to explore the data more specifically by gathering more descriptive statistics, especially about skew and kurtosis since those seem to be prevalent throughout the data set. Every interval category has a STDV larger than its mean, so it is likely worth the time to trim outliers from either end of this dataset so that it's more likely for analysis like hypothesis testing to be effective. I don't expect any of these categories to become normally distributed, but they almost definitely have extreme outliers that would warp any model that could be created based on the data as it is.

Additional data about Banned authors who posted Opinions could be a relevant part of a model as that is the main intersection with a small amount of values (less than 200).

Another piece of data that would be useful if it's accessible is the watch time for each video. Data about how long a user focused on each type of video might be an indicator the users impression of whether or not the video contains a claim and therefore could be a usable predictor in the model.

Data about how many followers each author had when the video was posted could also help in case the concept of "viralness" comes into the model.

[ ]: